

Day 1:

-----

- 1) Can we have AD group as job owner?
- 2) What does CheckDB with estimate only will give us?
- 3) When we need to keep DB in emergency mode? Do we need to do this for all DB corruptions?
- 4) If the corruption happened and you ran repair allow data loss how will ensure that there is no loss inside your database?
- 5) Can we have DB owner as AD group?
- 6) Will sp\_change\_users\_login report gives back orphan users at windows level?
- 7) How will you change the password for the distributor admin in replication?
- 8) What happens if we disable the clustered index?
- 9) Can we consider Log shipping as reporting solution?
- 10) In transactional replication will we get non clustered indexes by default to Subscribers?

Day 2:

-----

- 1) How to find out when CheckDB was last run?
- 2) What's the difference between Covering index and included columns?
- 3) What are incremental statistics?
- 4) What are the 3 modes inside sys.dm\_db\_index\_physical\_stats and can you run them during business hours for very large tables?
- 5) How will you track the percentage completion of index creation?
- 6) What is the difference between Availability Groups and Distributed Availability Groups?
- 7) When will you get Compute scalar operator inside Execution Plans?
- 8) What is cost threshold for parallelism and the things that prevent SQL to make use of Parallelism?
- 9) How do you know whether the stats were run with Full scan or sampling?
- 10) How will you synchronize the objects in Always ON that are outside the scope of the database?

Day 3:

-----

- 1) What's your understanding with Dbcc CheckDB messages that gets logged inside the error logs?
- 2) On Managed instance how many tempdb files would get created?

- 3) Do we need to create High Availability for Managed instance? If so how?
- 4) What's the difference between General Purpose and Business Critical Architectures?
- 5) What are different ways to connect to Azure Managed instance?
- 6) Have you ever made use of partitioned views in SQL Server?
- 7) How will you adjust the memory when we have multiple instances in case of standalone and clustered environments?
- 8) What's your take if total server memory is greater than target server memory?
- 9) What are few reasons that you think of for Always on Connection time outs?
- 10) I need to make a change for one of the columns from int to bigint it is of 1 TB how can you do this operation in seconds?

Day 4:

-----

- 1) If you are making use of Azure VM's what are the settings that you need to keep for Host caching with respect to data and log files?
- 2) If you're on-premises source version is anything less than 2016 and your target is PAAS (Azure Managed instance) and you need to move more than 50 databases with minimal Downtime what technology will you make use of?
- 3) If you're on-premises source version is 2016/2019 and your target is PAAS (Azure Managed instance) and the downtime is less than 5 minutes what technologies will you Make use of?
- 4) How will you access the data from one Azure SQL DB to another DB?
- 5) What's the difference between Bacpac and Dacpac?
- 6) How will Sargable expressions contribute to High CPU usage?
- 7) How will you find out if sqlservr.exe is the real contributor for high CPU usage if it happened in the past?
- 8) How will you find out Always on Failover History?
- 9) Will there be any default statistics on Secondary replicas? If so when would they get created and where will they be stored as the DB is read-only on secondary side?
- 10) Will there be any overhead on tempdb if we turn on readable secondaries?

Day 5:

-----

- 1) Which File group is the Database Log file created in?
- 2) Is it a good idea to keep "auto update statistics" enabled in SQL Server? Why? HOW?

- 3) Is disabling "parameter sniffing" a good performance tuning practice?
- 4) I have a 2 node AlwaysON setup. I received a request to create a login and grant permission on a Database which is part of that AlwaysON setup. Should i create "LOGIN" on either nodes or only primary node?
- 5) Can we create statistics under database tables or only the SQL server can do that?
- 6) What is the default value of the connection timeout property of "Availability Group" in SQL Server AlwaysON Availability Group configuration?
- 7) Should I create user on both nodes or only on primary node in Always ON?
- 8) Should I also grant permission on either nodes or just primary node?
- 9) Is it possible to get the actual execution plan from the cache?
- 10) What is subtree cost inside the execution plan?